

# **Phase 4 Report**

Distributed Systems, Spring 2022

Imani Muhammad-Graham

## Table of Contents

Introduction	.....	3
Design Overview	.....	3
Application	.....	7
Implementation	.....	7
Validation	.....	11
References	.....	13

## Introduction

### Project

The subject of interest of this topic is communication, featuring a “motivational quote” web application. The app allows users to post a quote that all other users can see. This report corresponds to Phase 4 of the Distributed Systems project.

### Phase 4

The motivation behind this Phase was to complete the implementation of the RAFT Consensus Algorithm that was started in Phase 3. Phase 4 incorporates the application designed in Phase 1&2 and the distributed cluster designed in Phase 3. The cluster would be tested and sampled with a controller. This phase entailed:

- Evolving Phase 2’s Application to a 5-node connected distributed system
- Incorporating Safe log replication into Phase 3’s design

The completion of Phase 2 required a demonstration of the following knowledge/skills:

- Network Design
- RAFT Consensus Algorithm
- Web Development

## Design Overview

The system’s design consists of multiple nodes working together to communicate with a client. Each node contains a server and database running in separate docker containers. The nodes communicate over a network instantiated by docker, and the server communicates with a client via web socket connections.

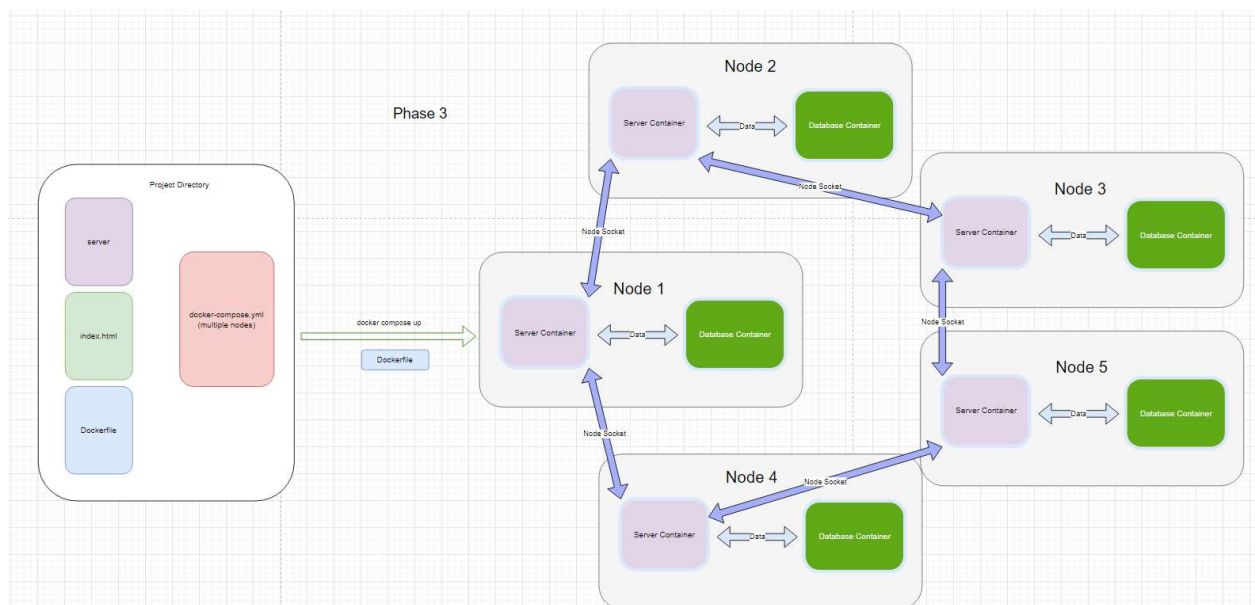


Figure 1: Model of system design (Part 1)

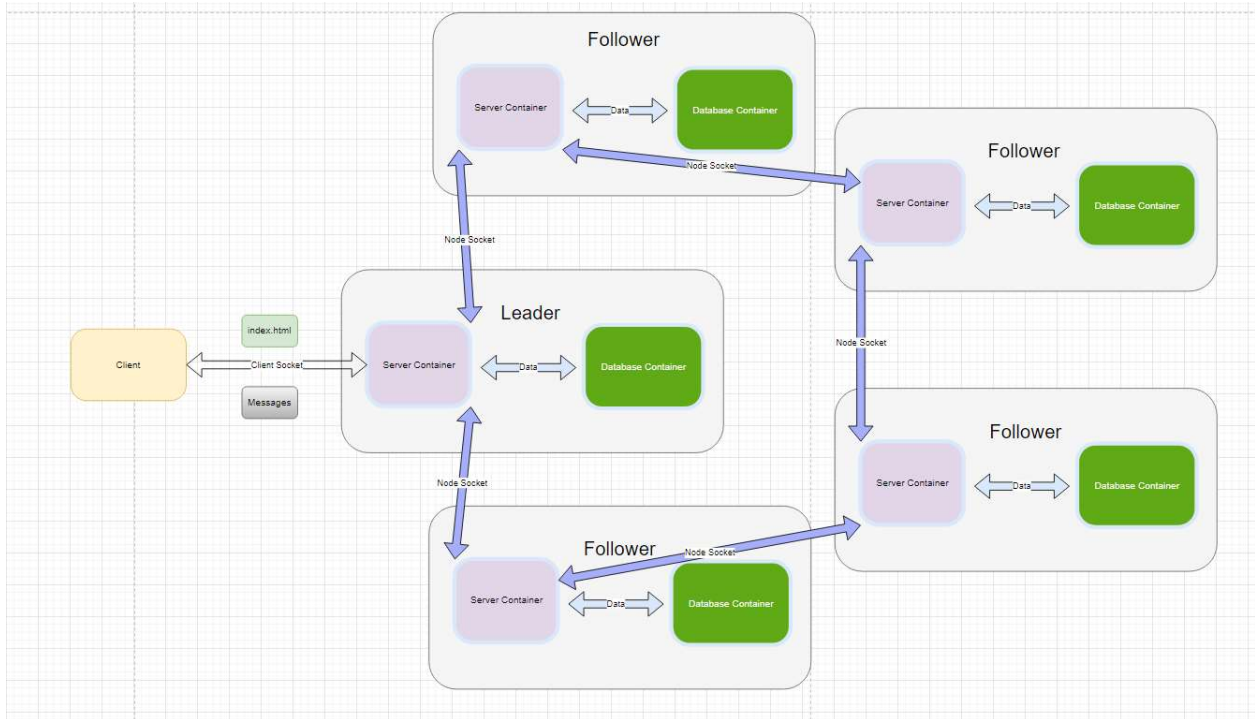


Figure 1: Model of system design (Part 2)

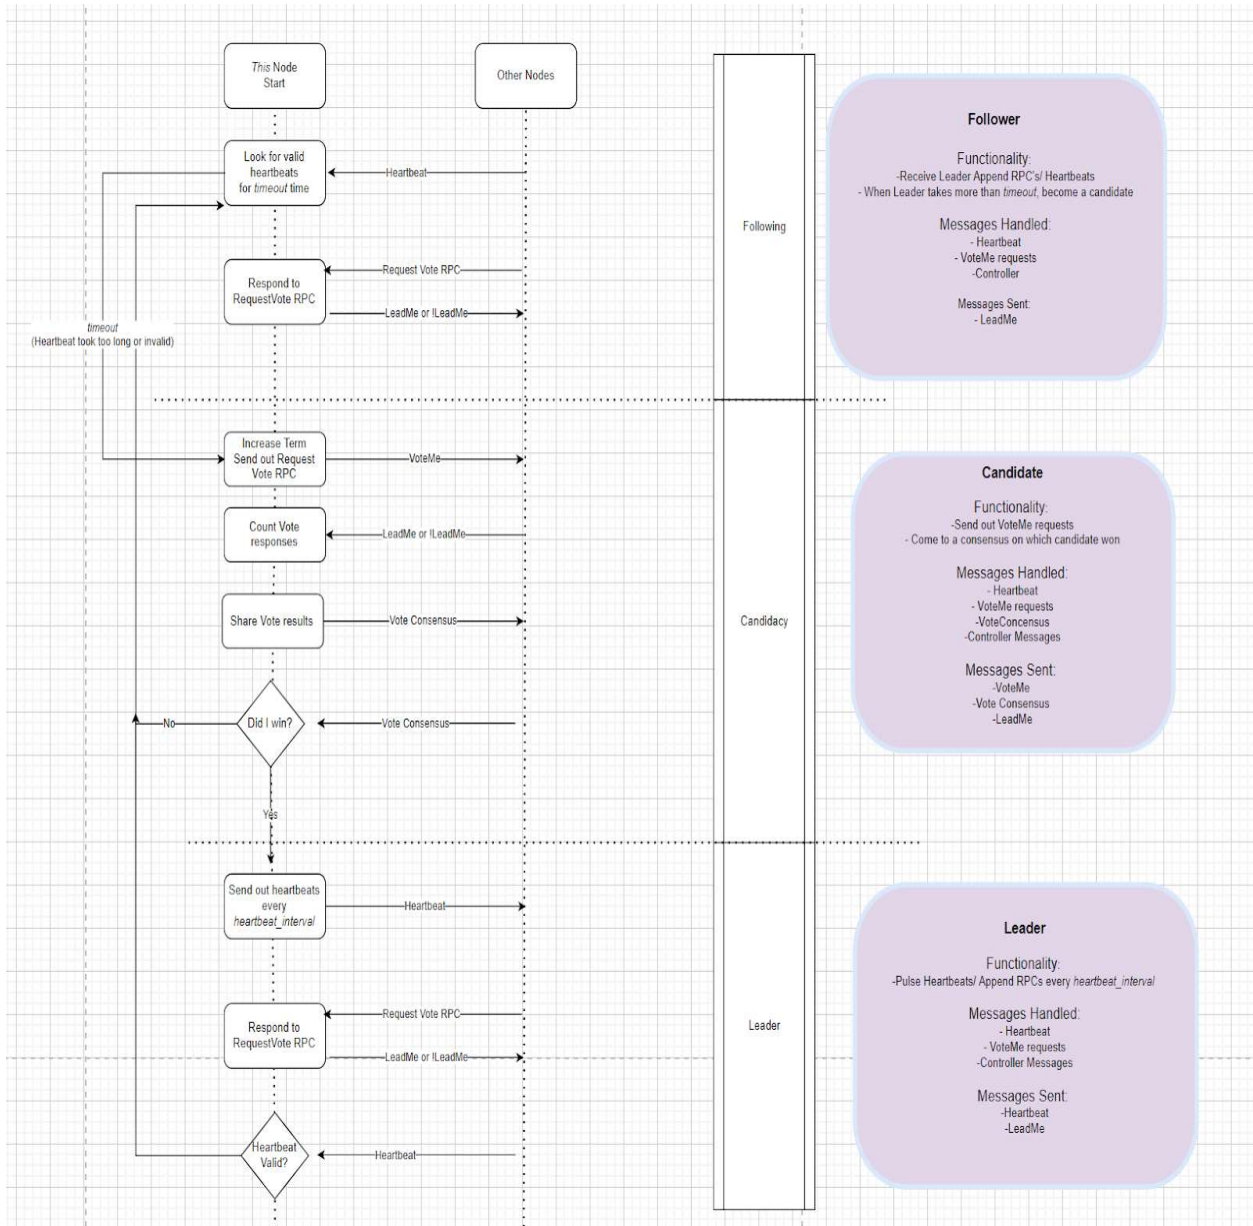


Figure 2: Functionality Overview (part 1)

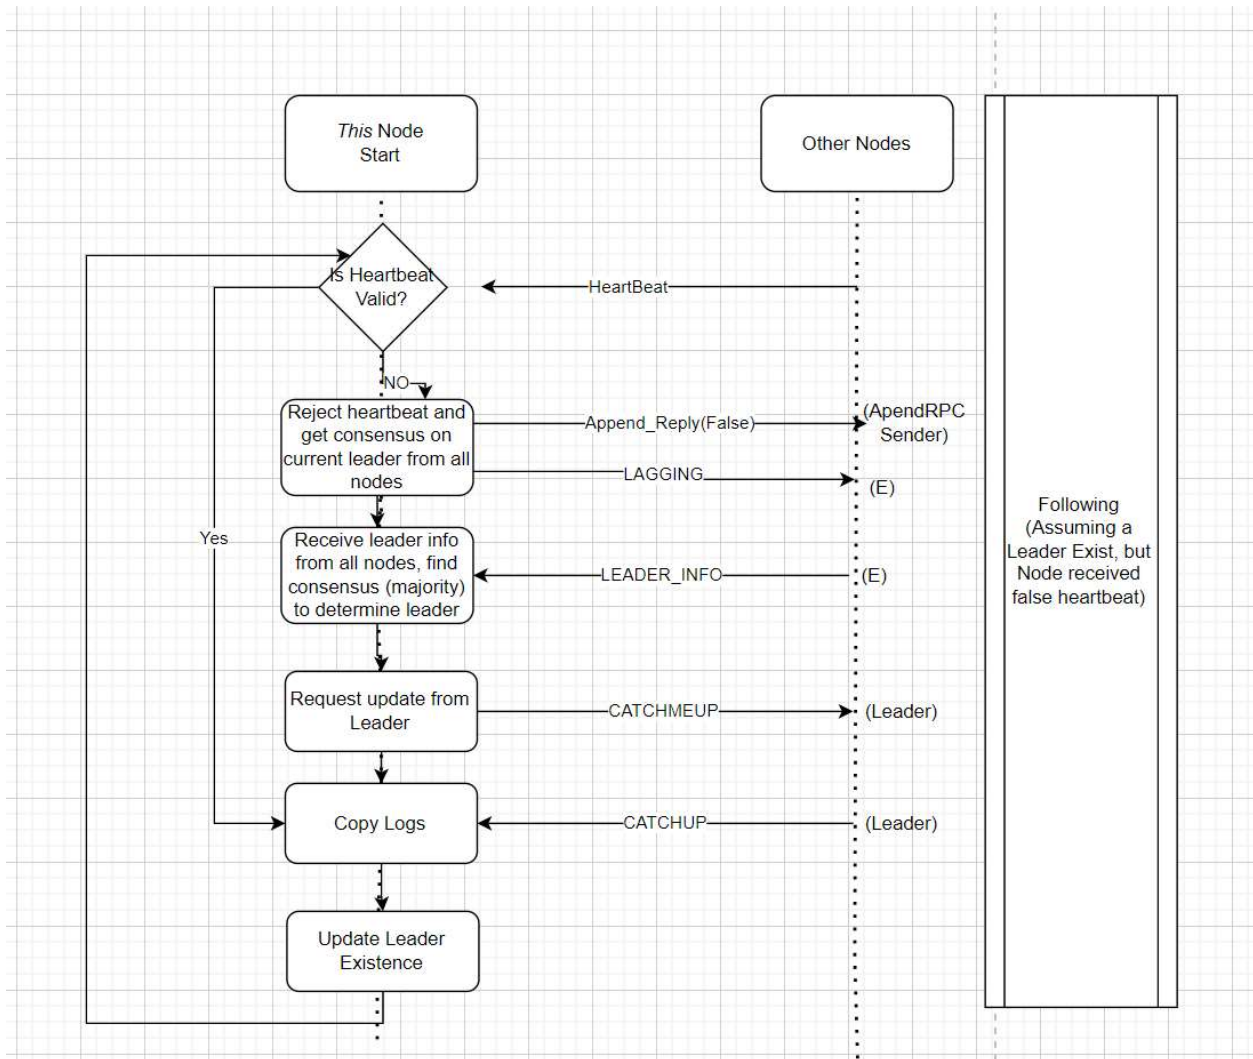


Figure 2: Functionality Overview (part 2)

This message architecture for Lagging nodes was chosen over 'leader's netindex'-via project description- because it allows for the Lagging node to get a consensus of who the leader is before trusting the leader that sent an AppendRPC. In the case that the Lagging nodes' own logs are corrupted, the logs won't be accepted from the leader, so CATCHUP contains all of the leaders' logs instead of one log at a time.

## Application

### The Application

This phase involved adding safe log replication to the consensus algorithm designed in Phase 3, as well as incorporating Phase 1's web application. The result is a simple RAFT-based cloud hosted web application.

The following tools were used to develop the application:

- Websockets
- Flask.io
- Javascript
- Python

### Docker

To deploy the app in its own container, 'docker-compose.yml' and 'Dockerfile' files were used. The docker-compose file sets up the networks, containers and the libraries and dependencies needed to run The Application, while the Dockerfile builds the container image.

## Implementation

### The Application

Safe log replication amongst the cluster and the web application were the two things implemented in Phase 4. For the Leader functionality, hosting of the web application and broadcasting AppendRPC's with Log information was added. Naturally, each nodes' message handler was upgraded to receive, judge, and append logs from the leader. Each state sent, and processed messages differently based on their roles, as seen in the figures below. Every node, no matter the state, had a listener running at all times. There is also a handler for messages from a Controller which is intended to change the state of the system with different commands including: SHUTDOWN, FOLLOW, TRYLEAD, STATUS, STORE, and RETRIEVE.

```

181     #node recieving command from controller
182     elif dm['sender_name'] == "Controller":
183         #UPDATE
184 >         if Request == 'STATUS':...
190         #turn node into a follower
191 >         elif Request == 'FOLLOW':...
194         #turn node into a candidate
195 >         elif Request == 'TRYLEAD':...
198         #have node play dead
199 >         elif Request == 'PLAYDEAD':...
202         #Add something to the log
203 >         elif Request == 'STORE':...
225         #Retrieve Logs
226 >         elif Request == 'RETRIEVE':...
239     #node recieving message from other node
240     elif node_info['state'] != 'd':
241         #follower recieving a candidate's vote request
242 >         if Request== "VOTEME":...
255         #candidate recieving a follower's vote respopnse
256 >         elif Request == "LEADME":...
259         #candidate recieving a follower's vote respopnse
260 >         elif Request == "!LEADME":...
262         #follower recieving a leader's heartbeat
263 >         elif Request == "HEARTBEAT":...
317         #Node recieving message from Lagging Node
318 >         elif Request == "LAGGING":...
325         #From node to Lagging Node
326 >         elif Request == "LEADER_INFO" and dm['recipient'] == name:...
329         #Lagging node to Leader
330 >         elif Request == "CATCHMEUP":...
337         #Leader trying to CATCHUP Lagging Node
338 >         elif Request == "CATCHUP":...
343         #Candidates checking for new leader
344 >         elif Request == "VOTECONCENSUS":...
347         #Append_Reply from followers to leader
348 >         elif Request == "Append_Reply":...
354         #Bad Message
355 >         else:...
357     #node is dead
358 >     else:...

```

Figure 3: Message Handler: The conditionals seen above in message\_handler() express the different types of messages handled from the controller or other nodes



```

263 elif Request == "HEARTBEAT":#verify heartbeat
264     #If node is lagging with respects to logs
265     if dm['prev_log_index'] > len(node_info['log']):#Send false success, broadcast Lagging ...
266     #Valid Heartbeat, caught up with logs
267     elif dm['term'] >= node_info['term'] and node_info['log'][dm['prev_log_index']]['term'] <= dm['prev_log_term']:#copy log...
268     #invalid heartbeat
269     else:#send failure RPC ...
270     #Node recieving message from Lagging Node
271     elif Request == "LAGGING":#send LEADER_INFO to lagging node...
272     #From node to Lagging Node
273     elif Request == "LEADER_INFO" and dm['recipient'] == name:#Get consensus from nodes on leader...
274     #Lagging node to Leader
275     elif Request == "CATCHMEUP": #Leader sending CATCHUP to logger...
276     #Leader trying to CATCHUP Lagging Node
277     elif Request == "CATCHUP": #Lagger recieving CATCHUP from leader...
278     #Candidates checking for new leader

```

Figure 4: Heartbeat handling

This follows figure 2's description of heartbeat sending and handling

## Docker

The dockerfiles below built the image for the containers

Phase_4 > RAFTBABY > Node > Dockerfile > FROM	Phase_4 > RAFTBABY > Controller > Dockerfile > FROM
1 FROM python:3.7-alpine3.14	1 FROM python:3.7-alpine3.14
2	2
3 RUN pip install flask	3 COPY . .
4 RUN pip install flask-socketio	4
5 # Stop the cache	5 EXPOSE 5555
6 ADD https://www.google.com /time.now	6
7	7 ENTRYPOINT ["python", "-u", "CNTRL.py"]
8 COPY . .	8
9	
10 EXPOSE 5555	
11	
12 ENTRYPOINT ["python", "-u", "node.py"]	

Figure 5: Dockerfile

Installs dependencies and runs server and controller containers. Google.com/time.now was to force new builds every time the container was run.

The docker-compose file below sets contexts and builds the server and controller containers for each node.

```

Phase_4 > RAFTBABY > 🐳 docker-compose.yml
6
7   version: "3.7"
8   services:
9 >   node1: ...
21 >  node2: ...
33 >   node3: ...
45 >   node4: ...
57   node5:
58     container_name: Node5
59     build: Node/.
60     environment:
61     - Port=5555
62     - app_name=Node5
63     - group=224.1.1.1
64     - tor1=5
65     - tor2=10
66     - heartrate=.25
67     - client_host=220.1.1.1
68     - client_port=5000
69
70   controller:
71     container_name: Controller
72     build: Controller/.
73     environment:
74     - Port=5555
75     - app_name=Controller
76     - group=224.1.1.1
77     stdin_open: true
78     depends_on:
79     node1:
80       condition: service_started
81 >   node2: ...
83 >   node3: ...
85 >   node4: ...
87 >   node5: ...

```

Figure 6: Docker-compose

Sets contexts, images, builds containers, defines port usage. The 5 nodes, and Controller were defined as such.

The command to run the application is:

`docker compose up --build`

## Validation

### Cluster up and running

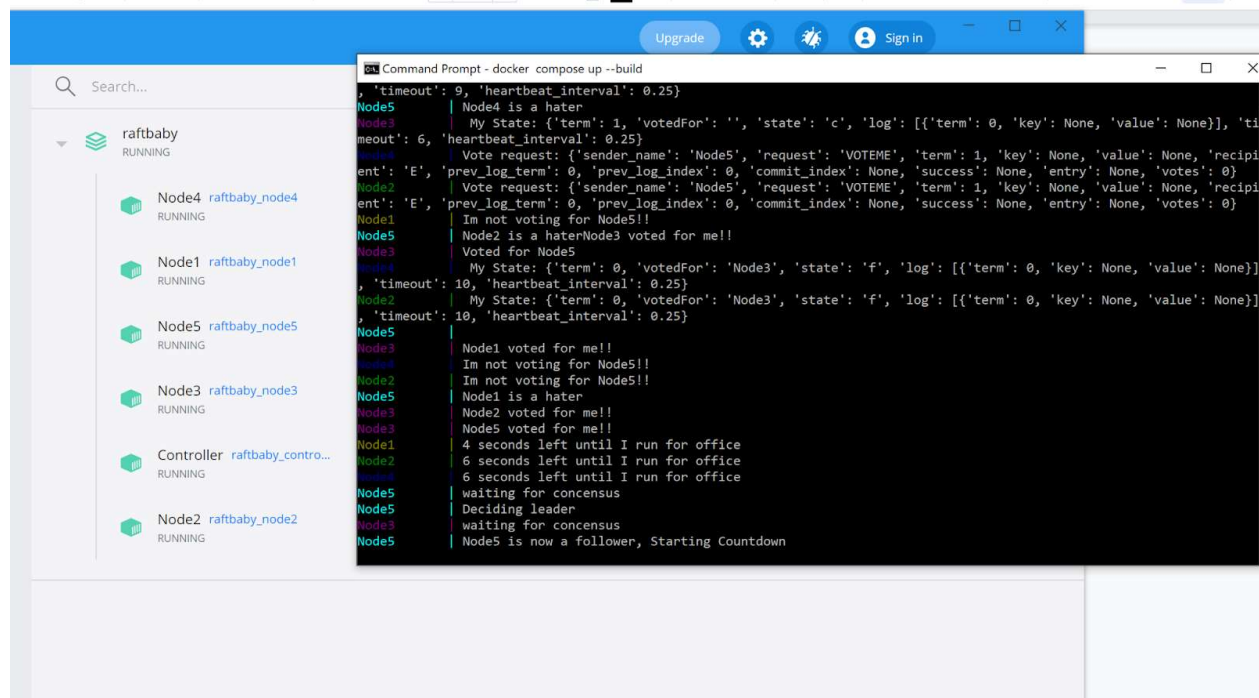


Figure 7: System running

Docker compose up –build was ran and created node containers. The nodes each started as a follower and were given a random timeout value, went into candidacy and are voting on the next leader.

```
203 elif Request == 'STORE':
204     if node_info['state'] == 'l':#store what controller wants
205         storethis = {'term':node_info['term'],'key':dm['key'],'value':dm['value']}
206         print(f"store the stuff controller wants: {storethis}")
207         node_info['log'].append(storethis)
208         #broadcasting appendrpc
209         msg_c['request'] = 'HEARTBEAT'
210         msg_c['recipient'] = 'E'
211         msg_c['term'] = node_info['term']
212         msg_c['prev_log_term'] = node_info['log'][len(node_info['log'])-2]['term']
213         msg_c['prev_log_term'] = len(node_info['log'])-2
214         msg_c['commit_index'] = len(node_info['log'])
215         msg_c['entry'] = storethis
216         send_message(msg_c,group,socket,port)
217         print(f"New log:{node_info['log']}")
218     else:#send leader_info to controller
219         print("Sending LEADER_INFO to controller")
220         msg_c['request']= 'LEADER_INFO'
221         msg_c['recipient'] = 'Controller'
222         msg_c['key'] = "LEADER"
223         msg_c['value'] = current_leader
224         send_message(msg_c,group,socket,port)
```

Figure 8: Log replication (Part 1: Node STORE functionality)

```

Controller raftbaby_controller
EXITED (137)

LOGS INSPECT STATS

Sending Request: {'sender_name': 'Controller', 'recipient': 'E', 'request': 'STORE', 'term': None, 'last_log': None, 'log_length': None, 'role': None, 'key': 'Cycle1', 'value': 'Chicken'}
Controller Received the following message: {'sender_name': 'Node1', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Controller Received the following message: {'sender_name': 'Node2', 'request': 'LEADER_INFO', 'term': 1, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': 0, 'prev_log_index': 0, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Controller Received the following message: {'sender_name': 'Node3', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Controller Received the following message: {'sender_name': 'Node4', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}

```

Figure 8: Log replication (Part 2: STORE)

```

Command Prompt

Node5 | Node2 Loves me :)Node1 Loves me :)
Node5 | Node4 Loves me :)
Node5 |
Node2 | All hail Node5
Node3 | All hail Node5
Node1 | All hail Node5
Node5 | Node4 Loves me :)
Node4 | All hail Node5
Node5 | Node2 Loves me :)Node1 Loves me :)
Node5 |
Node5 | Node3 Loves me :)
Controller | Sending Request: {'sender_name': 'Controller', 'recipient': 'E', 'request': 'STORE', 'term': None, 'last_log': None, 'log_length': None, 'role': None, 'key': 'Cycle1', 'value': 'Chicken'}
Controller | Controller Received the following message: {'sender_name': 'Node1', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Node3 | Sending LEADER_INFO to controller
Node4 | Sending LEADER_INFO to controller
Node2 | Sending LEADER_INFO to controller
Node1 | Sending LEADER_INFO to controller
Node5 | store the stuff controller wants: {'term': 1, 'key': 'Cycle1', 'value': 'Chicken'}
Controller | Controller Received the following message: {'sender_name': 'Node2', 'request': 'LEADER_INFO', 'term': 1, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': 0, 'prev_log_index': 0, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Controller | Controller Received the following message: {'sender_name': 'Node3', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Node3 | All hail Node5
Node4 | All hail Node5
Node2 | All hail Node5
Node1 | All hail Node5
Node5 | New log:[{'term': 0, 'key': None, 'value': None}, {'term': 1, 'key': 'Cycle1', 'value': 'Chicken'}]
Controller |
Node3 | All hail Node5
Node4 | All hail Node5
Node2 | All hail Node5
Node1 | All hail Node5
Node5 | Node4 Loves me :)
Controller | Controller Received the following message: {'sender_name': 'Node4', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Node5 | Node1 Loves me :)
Node5 | Node3 Loves me :)

```

Figure 8: Log replication (Part 3: Nodes react to STORE)

As you can see, the nodes have elected Node 5 as their leader. When the controller sends a STORE request, following nodes reply with Leader Information, while the Leader 'stores the stuff the controller wants'. Node 5's New Log print statement was added for testing.



```

Sending Request: {'sender_name': 'Controller', 'recipient': 'E', 'request': 'RETRIEVE', 'term': None, 'last_log': None, 'log_length': None,
'role': None, 'key': 'Cycle1', 'value': 'Chicken'}
Controller Received the following message: {'sender_name': 'Node1', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node
5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'vote
s': 0}
Controller Received the following message: {'sender_name': 'Node3', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node
5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'vote
s': 0}Controller Received the following message: {'sender_name': 'Node4', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value':
'Node5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None,
'votes': 0}

Controller Received the following message: {'sender_name': 'Node5', 'request': 'RETRIEVE', 'term': 1, 'key': 'COMMITTED_LOGS', 'value': [{'t
erm': 0, 'key': None, 'value': None}, {'term': 1, 'key': 'Cycle1', 'value': 'Chicken'}], 'recipient': 'Controller', 'prev_log_term': 1, 'pr
ev_log_index': 0, 'commit_index': 2, 'success': None, 'entry': None, 'votes': 0}
Controller Received the following message: {'sender_name': 'Node2', 'request': 'LEADER_INFO', 'term': 1, 'key': 'LEADER', 'value': 'Node5',
'recipient': 'Controller', 'prev_log_term': 0, 'prev_log_index': 0, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}
Sending Request: {'sender_name': 'Controller', 'recipient': 'E', 'request': 'STORE', 'term': None, 'last_log': None, 'log_length': None, 'r
ole': None, 'key': 'Cycle3', 'value': 'Chicken'}
Controller Received the following message: {'sender_name': 'Node3', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node
5', 'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'vote
s': 0}
Controller Received the following message: {'sender_name': 'Node2', 'request': 'LEADER_INFO', 'term': 1, 'key': 'LEADER', 'value': 'Node5',
'recipient': 'Controller', 'prev_log_term': 0, 'prev_log_index': 0, 'commit_index': None, 'success': 'true', 'entry': None, 'votes': 0}Con
troller Received the following message: {'sender_name': 'Node1', 'request': 'LEADER_INFO', 'term': None, 'key': 'LEADER', 'value': 'Node5',
'recipient': 'Controller', 'prev_log_term': None, 'prev_log_index': None, 'commit_index': None, 'success': 'true', 'entry': None, 'votes':
0}

```

Figure 8: Log Replication (Part 4: Retrieval of Stored Messages)

See video for a better visualization of the nodes' interactions, and what happens when you kill the leader.

## References:

[1]:

<http://www.steves-internet-guide.com/introduction-multicasting/#:~:text=Note%3A%20multicast%20uses%20UDP%20and,part%20of%20a%20multicast%20group>

[2]: [https://flask-socketio.readthedocs.io/en/latest/getting\\_started.html#initialization](https://flask-socketio.readthedocs.io/en/latest/getting_started.html#initialization)

[3]: [https://medium.com/@abhishekchaudhary\\_28536/building-apps-using-flask-socketio-and-javascript-socket-io-part-1-ae448768643](https://medium.com/@abhishekchaudhary_28536/building-apps-using-flask-socketio-and-javascript-socket-io-part-1-ae448768643)